


# Operating Systems

---



## Lecture 4: Scheduling

# Contact Information

Instructor: Engr. Shamila Nasreen

Assistant Professor

Department of Software Engineering

MUST

Email: [shamila.se@must.edu.com](mailto:shamila.se@must.edu.com)

Office hours: Wednesday: 8:30–10:00 AM

Thursday: 8:30–10:00 AM

# Objectives

- To introduce CPU scheduling, which is the basis for multiprogrammed operating systems.
- To describe various CPU-scheduling algorithms.
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system.
- To examine the scheduling algorithms of several operating systems.

# Today's Lecture Agenda

- Types of Processor Scheduling
  - Long Term Scheduling
  - Medium term Scheduling
  - Short Term Scheduling
- Scheduling Algorithms
- Short term scheduling Criteria

# Basic Concepts

- Maximum CPU utilization is obtained with multiprogramming
  - Several processes are kept in memory at one time
  - Every time a running process has to wait, another process can take over use of the CPU
- Scheduling of the CPU is fundamental to operating system design
- The CPU scheduler selects from among the processes in memory that are ready to execute and allocates the CPU to one of them

# Aims Of Scheduling

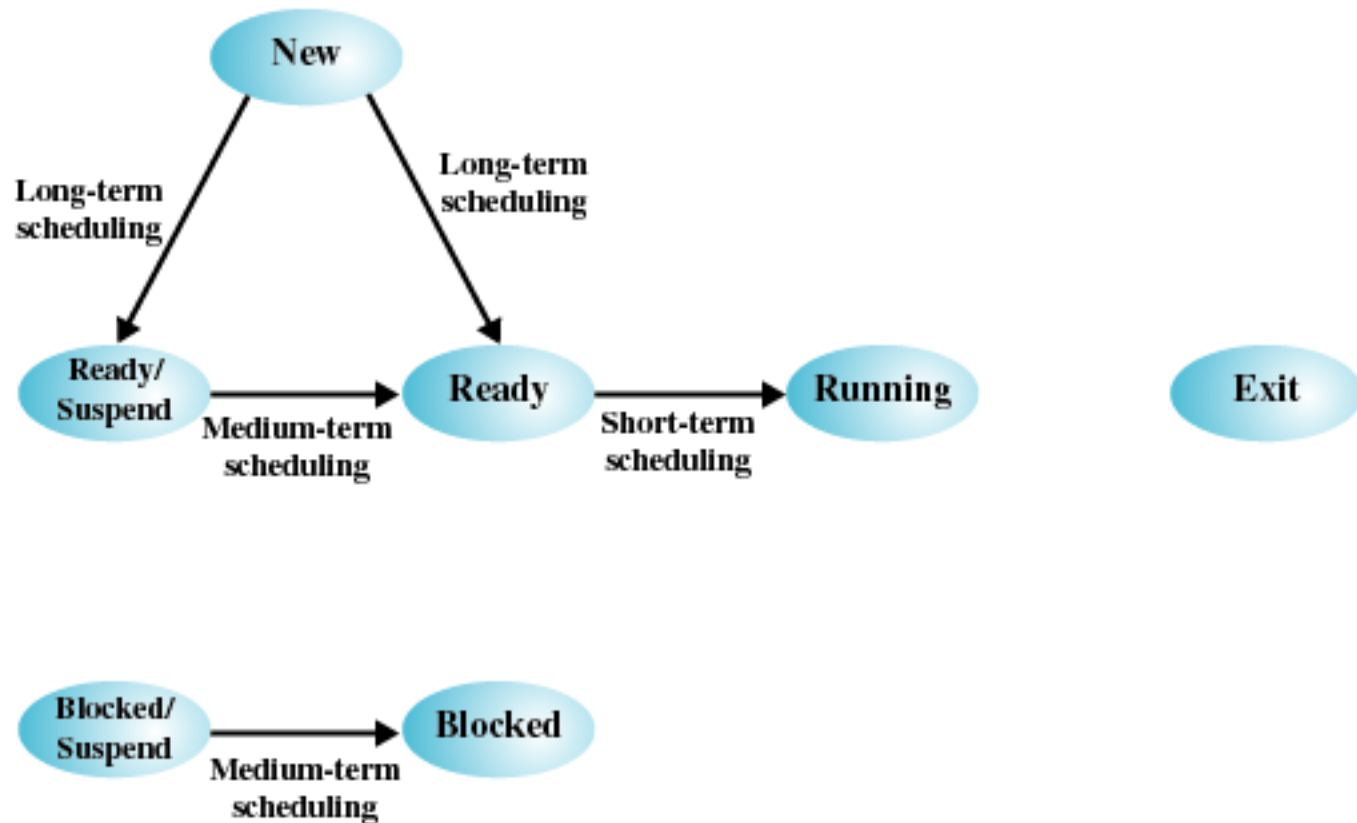
- The key to multiprogramming is scheduling.
- Aims:
  - Assign processes to be executed by the processor(s)
    -
  - Response time
  - Throughput
  - Processor efficiency
  - Fairness

# Types of Scheduling

**Table 9.1 Types of Scheduling**

<b>Long-term scheduling</b>	The decision to add to the pool of processes to be executed
<b>Medium-term scheduling</b>	The decision to add to the number of processes that are partially or fully in main memory
<b>Short-term scheduling</b>	The decision as to which available process will be executed by the processor
<b>I/O scheduling</b>	The decision as to which process's pending I/O request shall be handled by an available I/O device

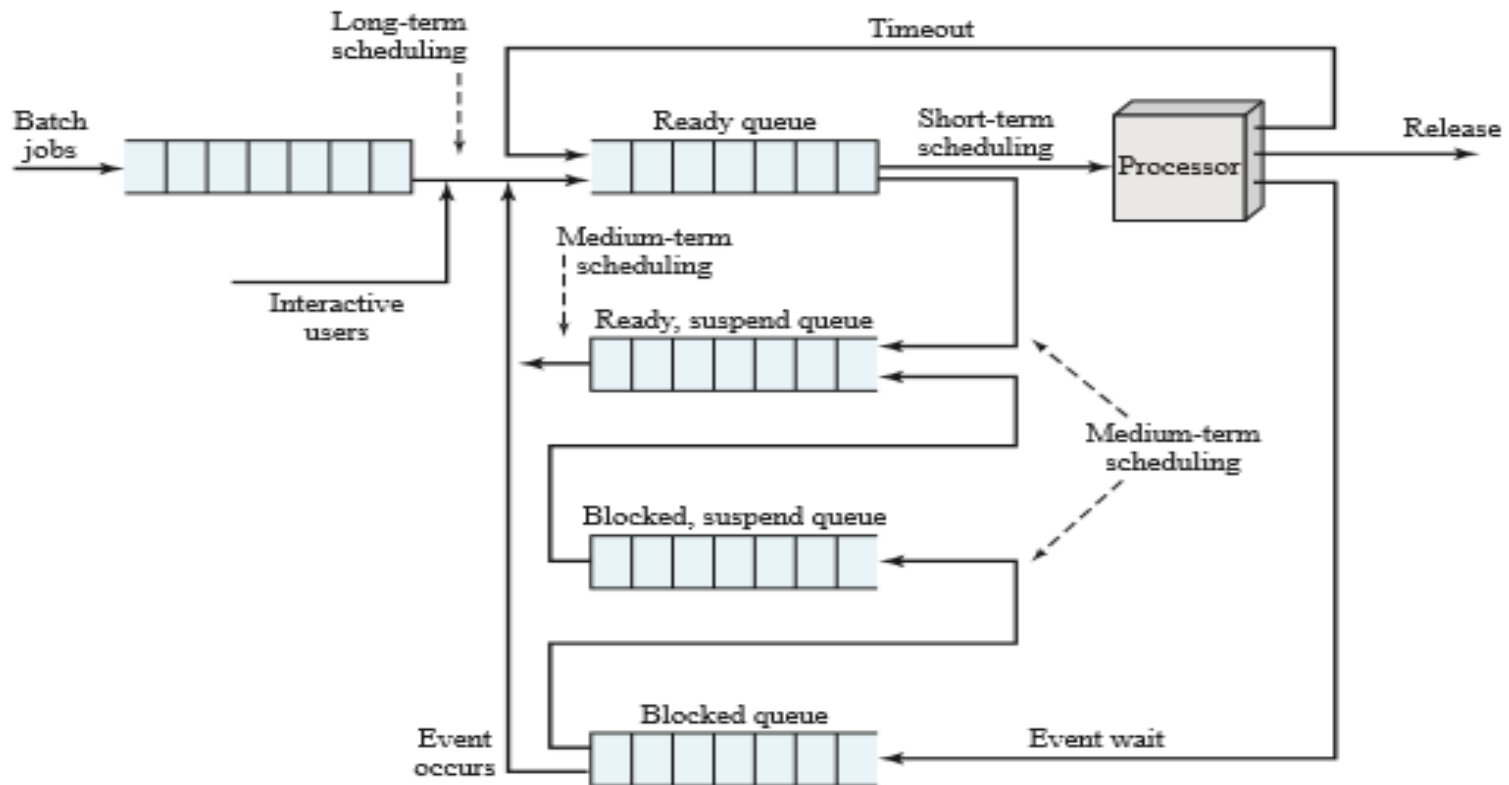
# Scheduling and Process State Transitions





# Medium term scheduling

- Part of the swapping function
- Based on the need to manage the degree of multiprogramming



**Figure 9.3** Queuing Diagram for Scheduling

# Short Term Scheduling

- In terms of Frequency of execution:
  - the **long-term scheduler** executes relatively infrequently
  - and makes the coarse-grained decision of whether or not to take on a new process and which one to take.
  - The **medium-term scheduler** is executed somewhat more frequently to make a swapping decision.
  - The **short-term scheduler** executes most frequently
  - and makes the fine-grained decision of which process to execute next.
- known as the dispatcher
- Dispatch latency – time it takes for the dispatcher to stop one process and start another running

# Short Term Scheduling

Points in time when OS may perform process scheduling:

- When a process blocks waiting for an event.
  - Perform a system call.
  - I/O request
  - Wait() invocation
- When an interrupt happens.
  - Clock interrupt.
  - I/O interrupt.
- When a process switches from waiting state to ready state
  - I/O completion
- When a process ends.

# Short Term Scheduling Criteria

- Generally, a set of criteria is established against which various scheduling policies may be evaluated.
- **User Oriented Criteria:**
  - User oriented criteria relate to the behavior of the system as perceived by the individual user or process.
  - An example is **response time** in an interactive system.
  - Response time is the elapsed time between the submission of a request until the response begins to appear as output.
  - This quantity is visible to the user and is naturally of interest to the user.
  - We would like a scheduling policy that provides “**good**” service to various users.
  - In the case of response time, a threshold may be defined, say 2 seconds.
  - Then a goal of the scheduling mechanism should be to maximize the number of users who experience an average response time of 2 seconds or less.

# Short Term Scheduling Criteria

- **System Oriented Criteria:**

- focus is on effective and efficient utilization of the processor.
- An example is **throughput**, which is the rate at which processes are completed.
- This is certainly a worthwhile measure of system performance and one that we would like to maximize.
- However, it focuses on system performance rather than service provided to the user.

# Decision Mode of Scheduling

- Nonpreemptive
  - Once a process is in the running state, it will continue until it terminates or blocks itself for I/O
- Preemptive
  - Currently running process may be interrupted and moved to the Ready state by the operating system
  - Allows for better service since any one process cannot monopolize the processor for very long

# Scheduling Criteria

- Different CPU scheduling algorithms have different properties
- The choice of a particular algorithm may favor one class of processes over another
- In choosing which algorithm to use, the properties of the various algorithms should be considered
- Criteria for comparing CPU scheduling algorithms may include the following
  - **CPU utilization** – percent of time that the CPU is busy executing a process
  - **Throughput** – number of processes that are completed per time unit
  - **Response time** – amount of time it takes from when a request was submitted until the first response occurs (but not the time it takes to output the entire response)
  - **Waiting time** – the amount of time before a process starts after first entering the ready queue (or the sum of the amount of time a process has spent waiting in the ready queue)
  - **Turnaround time( $T_q$ )** – amount of time to execute a particular process from the time of submission through the time of completion
    - Overall time a process is in system.
    - $T_q = T_f - T_i$
    - $T_f$ : Finalization time.
    - $T_i$ : Initiation time.

# Optimization Criteria

- It is desirable to
  - Maximize CPU utilization
  - Maximize throughput
  - Minimize turnaround time
  - Minimize start time
  - Minimize waiting time
  - Minimize response time
- In most cases, we strive to optimize the average measure of each metric



# First come, First serve


- Simplest scheduling algorithm:
  - Run jobs in order that they arrive
- Non-preemptive
  - A Process keeps CPU until done(terminate) or requesting I/O
- Advantage:
  - Simplicity

# First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- With FCFS, the process that requests the CPU first is allocated the CPU first
- Case #1: Suppose that the processes arrive in the order:  $P_1$  ,  $P_2$  ,  $P_3$


The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27) / 3 = 17$
- Average turn-around time:  $(24 + 27 + 30) / 3 = 27$

# First come, First serve


- Disadvantage
  - Wait time depends on arrival order
  - Unfair to later jobs
  - (worst case: long job arrives first)
  - It is a troublesome algorithm for time-sharing systems
- Three jobs (times: A=100, B=1, C=2) arrive in the order A, B, C



$$\begin{aligned}\text{Average} \\ \text{Waiting Time} &= (0 + 100 + 101) / 3 \\ &= 67\end{aligned}$$

# First come, First serve

- Now if they arrive in the order B, C, A



$$\begin{aligned}\text{Average} \\ \text{Waiting Time} &= (0 + 1 + 3) / 3 \\ &= 1.33\end{aligned}$$

# FCFS Convoy effect

- A CPU bound job will hold CPU until
  - Terminates
  - Or it causes an I/O burst
    - ⌞ Rare occurrence, since the process is CPU-bound
- Long periods where no I/O requests issued, and CPU held
- Result:
  - Poor I/O device utilization
- Penalizes short processes: *Convoy effect* short process behind long process

# Activity #1

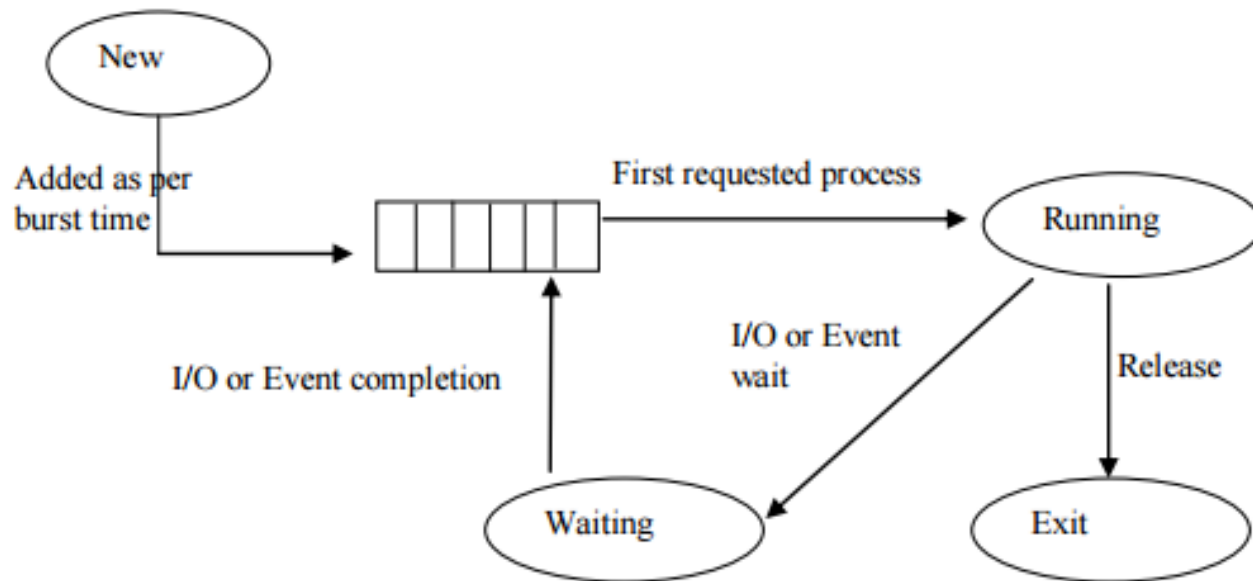
Process	Arrival	Service	Initialize	End	Wait
A	0	3			
B	2	6			
C	4	4			
D	6	5			
E	8	2			

Average waiting time?  
Average turn around time?

# Shortest-Job-First (SJF) Scheduling

- The SJF algorithm associates with each process the length of its next CPU burst
- When the CPU becomes available, it is assigned to the process that has the smallest next CPU burst (in the case of matching bursts, FCFS is used)
- Two schemes:
  - **Nonpreemptive** – once the CPU is given to the process, it cannot be preempted until it completes its CPU burst
  - **Preemptive** – if a new process arrives with a CPU burst length less than the remaining time of the current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)

# Shortest-Job-First (SJF) Scheduling






## Example #1: Non-Preemptive SJF (simultaneous arrival)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	6
$P_2$	0.0	4
$P_3$	0.0	1
$P_4$	0.0	5

- SJF (non-preemptive, simultaneous arrival)




- Average waiting time =  $(0 + 1 + 5 + 10)/4 = 4$
- Average turn-around time =  $(1 + 5 + 10 + 16)/4 = 8$

## Example #2: Non- Pre-emptive SJF (varied arrival times)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (non- preemptive, varied arrival times)



- Average waiting time  

$$= ( (0 - 0) + (8 - 2) + (7 - 4) + (12 - 5) ) / 4$$

$$= (0 + 6 + 3 + 7) / 4 = 4$$
- Average turn-around time:  

$$= ( (7 - 0) + (12 - 2) + (8 - 4) + (16 - 5) ) / 4$$

$$= (7 + 10 + 4 + 11) / 4 = 8$$

Waiting time : sum of time that a process has spent waiting in the ready queue

# Shortest Job First(SJF)

- Non-preemptive algorithm.
  - Selects shortest job.
- It can only be applied if duration of each job is known beforehand.
- SJF is optimal – gives minimum average waiting time for a given set of processes
  - The difficulty is knowing the length of the next CPU request
- Starvation possibility:
  - If short jobs are continuously arriving, longer jobs never are in position to be executed.


## Activity #2

Process	Arrival	Service	Initialize	End	Wait
A	0	3			
B	2	6			
C	4	4			
D	6	5			
E	8	2			

Average waiting time?  
Average turn around time?

## Activity #2


Process	Arrival	Service	Initialize	End	Wait
A	0	3			
B	2	6			
C	4	4			
D	6	5			
E	8	2			



# Shortest-remaining-time-first(SRTF)

- Preemptive version of shortest Job next policy

Process	Arrival Time	Burst Time	Remaining Time		
$P_1$	0.0	7	5	5	5
$P_2$	2.0	4	2	0	0
$P_3$	4.0	1	0	0	0
$P_4$	5.0	4	4	4	0



- Average waiting time  

$$= ( [(0 - 0) + (11 - 2)] + [(2 - 2) + (5 - 4)] + (4 - 4) + (7 - 5) ) / 4$$

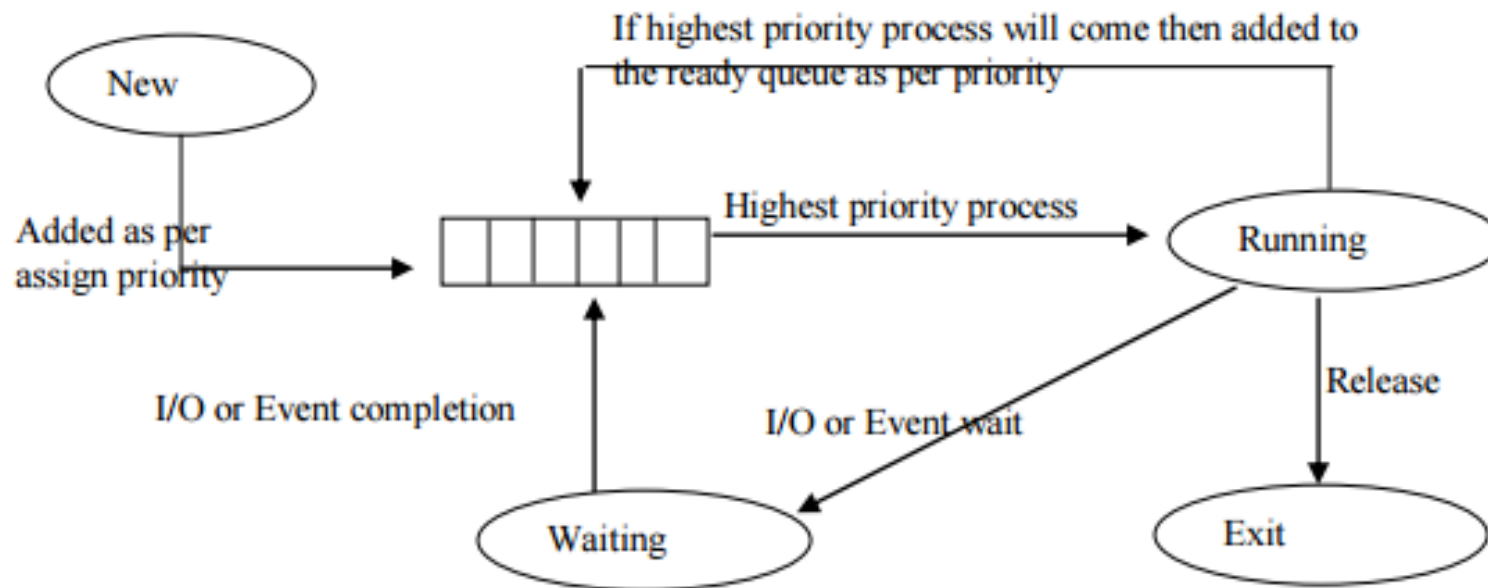
$$= 9 + 1 + 0 + 2 / 4$$

$$= 3$$
- Average turn-around time = ??

# Priority Scheduling

- The OS assigns a fixed priority rank to every process.
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
- The scheduler arranges the processes in the ready queue in order of their priority.
- Lower priority processes get interrupted by incoming higher priority processes.
- Characteristics
  - Starvation can happen to the low priority process.
  - The waiting time gradually increases for the equal priority processes .
  - Higher priority processes have smaller waiting time and response time.

# Priority Scheduling






# Priority Scheduling

- The SJF algorithm is a special case of the general priority scheduling algorithm
  - priority is the predicted next CPU burst time
- Priorities can be static or dynamic or both
- Among the Processes of equal priority
  - Round robin
  - FCFS

# Priority Scheduling

Priority is assigned for each process as follows.

Process ID	Burst Time(ms)	Priority
P0	12	3
P1	2	1
P2	3	3
P3	2	4
P4	6	2



**Gantt Chart for priority Scheduling**

**Waiting Time:**

P0=8, P1=0, P2=20, P3=23, P4=2

**Avg Waiting Time:**  $8+0+20+23+2/5 = 10.6$

**Avg Turnaround Time:**  $20+2+23+25+8/5 = 15.6$

# Priority scheduling

- Priority scheduling can be Preemptive or Non-Preemptive
- When a process arrives and enters the Ready Queue
- Its priority is compared with the currently Running Process
- If Higher
  - Preemptive Scheduling
    - ⌞ Run the New Thread
  - Non-Preemptive Scheduling
    - ⌞ Continue running the Current process.

# Priority scheduling

- High priority always runs over low priority.
- **Starvation**
  - A low Priority process may indefinitely wait for the CPU
- Solution: **Aging**
  - Gradually increase the Priority of processes that wait in the system for a long time.

# Round robin (RR)

- Solution to job monopolizing CPU? Interrupt it.
  - Run job for some “time slice,”
  - When time is up, or it blocks
  - It moves to back of a FIFO queue
- Advantage:
  - Fair allocation of CPU across jobs
  - Low average waiting time when job lengths vary

# Round robin (RR)

- In the round robin algorithm, each process gets a small unit of CPU time (a *time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Performance of the round robin algorithm
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow q$  must be greater than the context switch time; otherwise, the overhead is too high
- One rule of thumb is that 80% of the CPU bursts should be shorter than the time quantum

# Round Robin's Disadvantage

- Good for Varying sized jobs
- But what about same-sized jobs?
- Assume 2 jobs of time =100 each:



- Avg completion time?
- $(200 + 200) / 2 = 200$
- How does this compare with FCFS for same two jobs?
- $(100 + 200) / 2 = 150$

# RR Time slice tradeoffs

- Performance depends on length of the timeslice
- Context switching isn't a free operation.
- If timeslice time is set too high (attempting to amortize context switch cost)
  - You get FCFS.
  - i.e. Processes will finish or block before their slice is up anyway
- If it's set too low you're spending all of your time context switching between threads.



# Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
----------------	-------------------


$P_1$	53
-------	----

$P_2$	17
-------	----

$P_3$	68
-------	----

$P_4$	24
-------	----

- The Gantt chart is:



0    20    37    57    77    97    117    121    134    154    162

- Typically, higher average turnaround than SJF, but better *response time*
- Average waiting time  


$$= ( [(0 - 0) + (77 - 20) + (121 - 97)] + (20 - 0) + [(37 - 0) + (97 - 57) + (134 - 117)] + [(57 - 0) + (117 - 77)] ) / 4$$

$$= (0 + 57 + 24) + 20 + (37 + 40 + 17) + (57 + 40) / 4$$

$$= (81 + 20 + 94 + 97) / 4$$

$$= 292 / 4 = 73$$
- Average turn-around time =  $(134 + 37 + 162 + 121) / 4 = 113.5$

# Time Quantum and Context Switches




## Activity #3 (q=1)

Process	Arrival	Service	Initialize	End	Wait
A	0	3			
B	2	6			
C	4	4			
D	6	5			
E	8	2			

Average waiting time?  
Average turn around time?

## Activity #3 (q=1)

Process	Arrival	Service	Initialize	End	Wait
A	0	3			
B	2	6			
C	4	4			
D	6	5			
E	8	2			



# Highest Response Ratio Next (HRRN)


- Choose next process with the greatest ratio
- time spent waiting + expected service time
- $R = w + s/s$
- $R$  = response ratio
- $W$  = time spent waiting for processor
- $S$  = expected service time
- (HRRN) scheduling is a non-preemptive discipline
- Jobs gain higher priority the longer they wait, which prevents indefinite postponement (process starvation).

# Highest Response Ratio Next (HRRN)

Table 9.4 Process Scheduling Example

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

Highest Response  
Ratio Next (HRRN)



# Multi-level Queue Scheduling

- Multi-level queue scheduling is used when processes can be classified into groups
- For example, **foreground** (interactive) processes and **background** (batch) processes
  - The two types of processes have different response-time requirements and so may have different scheduling needs
  - Also, foreground processes may have priority (externally defined) over background processes
- A multi-level queue scheduling algorithm **partitions the ready queue** into several separate queues
- The processes are permanently assigned to one queue, generally based on some property of the process such as memory size, process priority, or process type
- Each queue has its own scheduling algorithm
  - The foreground queue might be scheduled using an RR algorithm
  - The background queue might be scheduled using an FCFS algorithm
- In addition, there needs to be scheduling among the queues, which is commonly implemented as fixed-priority pre-emptive scheduling
  - The foreground queue may have absolute priority over the background queue

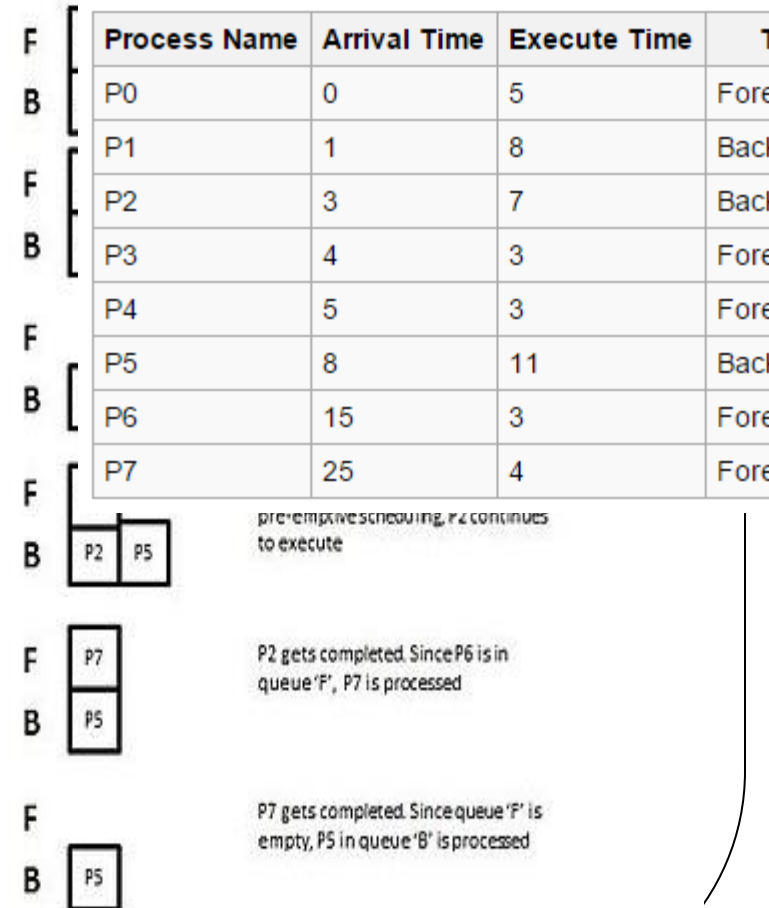
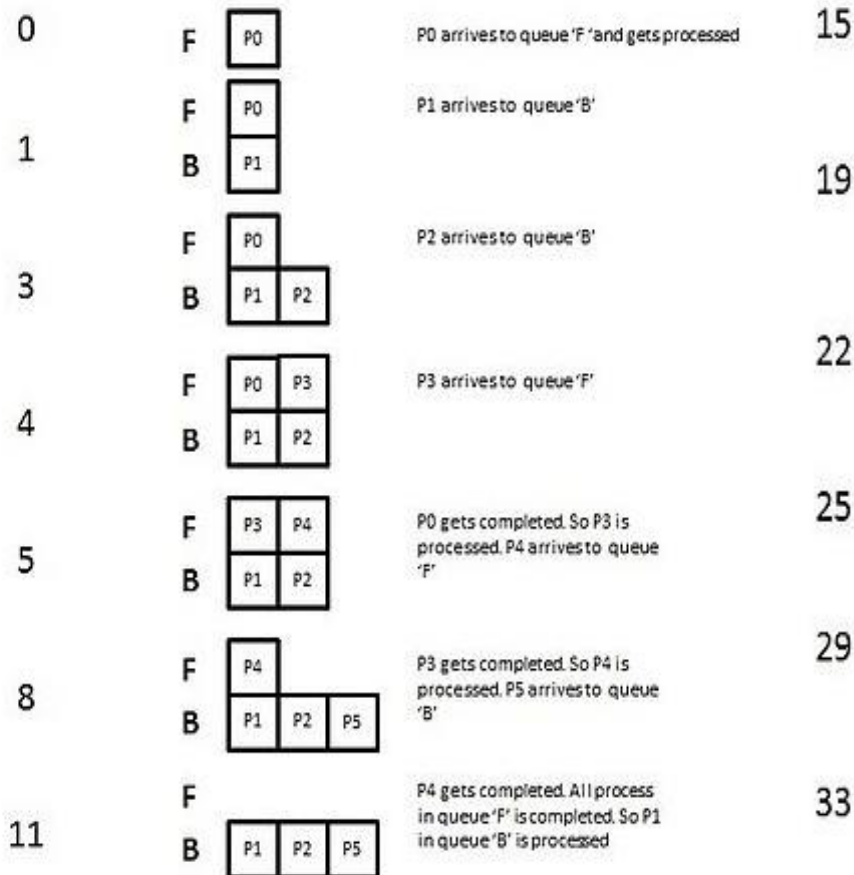
# Multi-level Queue Scheduling

Process Name	Arrival Time	Execute Time	Type
P0	0	5	Foreground
P1	1	8	Background
P2	3	7	Background
P3	4	3	Foreground
P4	5	3	Foreground
P5	8	11	Background
P6	15	3	Foreground
P7	25	4	Foreground



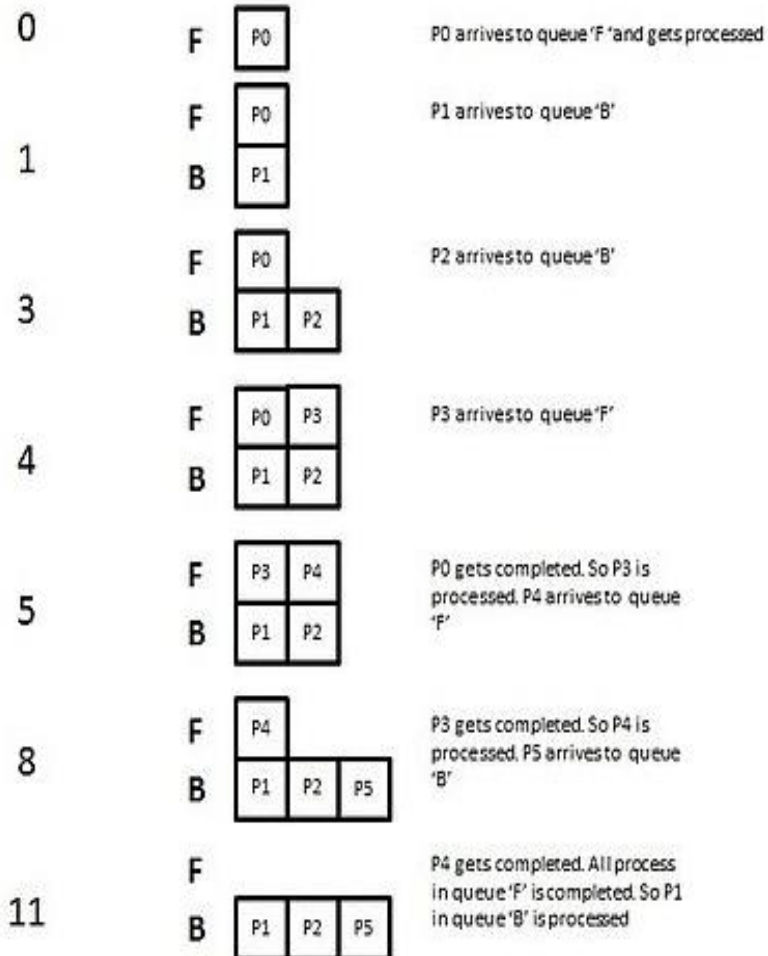
# Multi-level Queue Scheduling<sub>(non Preemptive)</sub>

## Execute Time Non Pre-emptive Process Scheduling



# Multi-level Queue Scheduling

## Execute Time Pre-emptive Process Scheduling



15

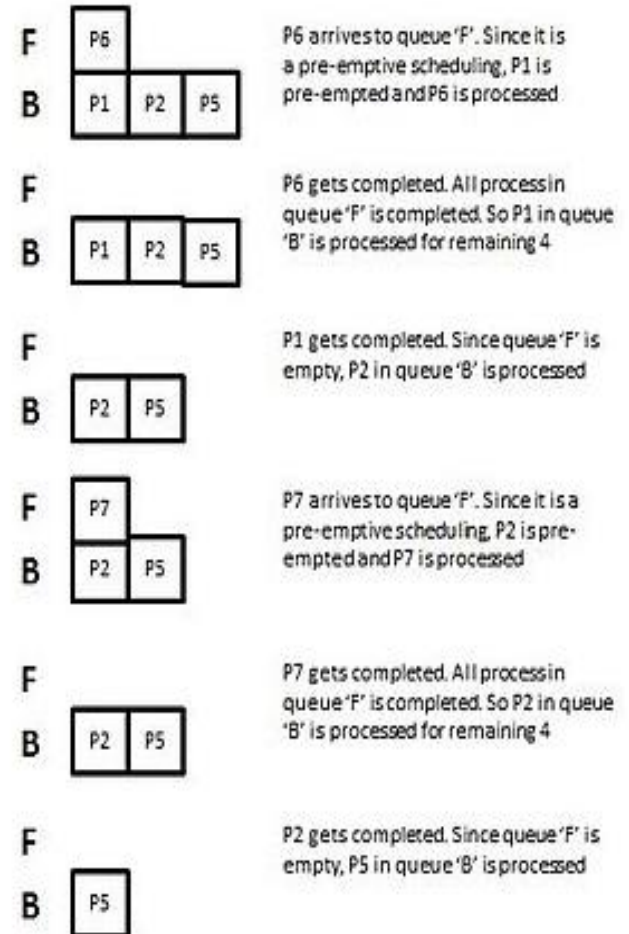
18

22

25

29

33




# Multilevel Feedback Queue Scheduling

- In multi-level feedback queue scheduling, a process can move between the various queues; aging can be implemented this way
- A multilevel-feedback-queue scheduler is defined by the following parameters:
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to determine when to promote a process
  - Method used to determine when to demote a process
  - Method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue


- Scheduling
  - A new job enters queue  $Q_0$  (RR) and is placed at the end. When it gains the CPU, the job receives 8 milliseconds. If it does not finish in 8 milliseconds, the job is moved to the end of queue  $Q_1$ .
  - A  $Q_1$  (RR) job receives 16 milliseconds. If it still does not complete, it is pre-empted and moved to queue  $Q_2$  (FCFS).



# Example of Multilevel Feedback Queue

- At the base level queue the processes circulate in round robin fashion until they complete and leave the system.
  - Processes in the base level queue can also be scheduled on a FCFS basis.
- Optionally, if a process blocks for I/O, it is 'promoted' one level, and placed at the end of the next-higher queue. This allows I/O bound processes to be favored by the scheduler and allows processes to 'escape' the base level queue.
- Only if the highest level queue has become empty will the scheduler take up a process from the next lower level queue.
  - The same policy is implemented for picking up in the subsequent lower level queues.
    - τ Meanwhile, if a process comes into any of the higher level queues, it will preempt a process in the lower level queue.

# MLBQ



**Figure 9.10** Feedback Scheduling

# Example of Multilevel Feedback Queue

Process	Burst Time
---------	------------

P1	30
----	----

P2	20
----	----

P3	10
----	----

The system has three RR queues with the following time slices:  
1 for queue1, 2 for queue 2, and 4 for queue3.

Draw the Gantt Chart

# Lottery scheduling: random simplicity

- **Lottery scheduling! Very simple idea:**
  - give each process some number of lottery tickets
  - On each scheduling event, randomly pick ticket
  - run winning process
- **How to use?**
  - Approximate priority: low-priority, give few tickets, high-priority give many
  - Approximate SJF: give short jobs more tickets, long jobs fewer. Key: If job has at least 1, will not starve
  - Degrades gracefully as load changes. Adding or deleting a job affects all jobs proportionately, independent of the number of tickets a job has.



# Lottery scheduling

- Example: give all jobs  $1/n$  of cpu?

- 4 jobs, 1 ticket each



- each gets (on average) 25% of CPU.

- Delete one job:



- automatically adjusts to 33% of CPU!

# Lottery scheduling

- Short jobs get 10 tickets, long jobs get 1 ticket each.

# short jobs/ # long jobs	% of CPU each short job gets	% of CPU each long job gets
1/1	91%	9%
0/2		
2/0		
10/1		
1/10		

# Lottery scheduling

- Short jobs get 10 tickets, long jobs get 1 ticket each.

# short jobs/ # long jobs	% of CPU each short job gets	% of CPU each long job gets
1/1	91% (10/11)	9% (1/11)
0/2		50% (1/2)
2/0	50% (10/20)	
10/1	10% (10/101)	< 1% (1/101)
1/10	50% (10/20)	5% (1/20)

# Example

- Three threads
  - A has 5 tickets
  - B has 3 tickets
  - C has 2 tickets
- If all compete for the resource
  - B has 30% chance of being selected
- If only B and C compete
  - B has 60% chance of being selected
- Reading Assignment: Is Lottery scheduling is ***starvation-free?***
-



## **5.4 Multiple-Processor Scheduling**

# Multiple-Processor Scheduling

- If multiple CPUs are available, load sharing among them becomes possible; the scheduling problem becomes more complex
- We concentrate in this discussion on systems in which the processors are identical (homogeneous) in terms of their functionality
  - We can use any available processor to run any process in the queue
- Two approaches: **Asymmetric** processing and **symmetric** processing (see next slide)

# Multiple-Processor Scheduling


- *Asymmetric multiprocessing (ASMP)*
  - One processor handles all scheduling decisions, I/O processing, and other system activities
  - The other processors execute only user code
  - Because only one processor accesses the system data structures, the need for data sharing is reduced
- Symmetric multiprocessing (SMP)
  - Each processor schedules itself
  - All processes may be in a common ready queue or each processor may have its own ready queue
  - Either way, each processor examines the ready queue and selects a process to execute
  - Efficient use of the CPUs requires load balancing to keep the workload evenly distributed
    - ⌞ In a **Push** migration approach, a specific task regularly checks the processor loads and redistributes the waiting processes as needed
    - ⌞ In a **Pull** migration approach, an idle processor pulls a waiting job from the queue of a busy processor
  - Virtually all modern operating systems support SMP, including Windows XP, Solaris, Linux, and Mac OS X

# Symmetric Multithreading

- Symmetric multiprocessing systems allow several threads to run concurrently by providing multiple physical processors
- An alternative approach is to provide multiple **logical** rather than **physical** processors
- Such a strategy is known as symmetric multithreading (SMT)
  - This is also known as hyperthreading technology
- The idea behind SMT is to create multiple logical processors on the same physical processor
  - This presents a view of several logical processors to the operating system, even on a system with a single physical processor
  - Each logical processor has its own **architecture state**, which includes general-purpose and machine-state registers
  - Each logical processor is responsible for its own interrupt handling
  - However, each logical processor shares the resources of its physical processor, such as cache memory and buses
- SMT is a feature provided in the hardware, not the software
  - The hardware must provide the representation of the architecture state for each logical processor, as well as interrupt handling (see next slide)



# A typical SMT architecture



SMT = Symmetric Multi-threading